



SRM
INSTITUTE OF SCIENCE & TECHNOLOGY
(Deemed to be University u/s 3 of UGC Act, 1956)

Heart Disease Prediction End-To-End data pipeline using AWS

By:

Name: Asmit Srivastava, Reg. No: [RA2512033010006](#)

Name: Biswadip Ghosh, Reg. No: [RA2512033010007](#)

Name: P.Bhuvanesh , Reg. No: [RA2512033010033](#)

Abstract

This research presents an automated end-to-end heart disease prediction system built on AWS cloud services, integrating Amazon S3, AWS Glue, Amazon EC2, and Amazon SageMaker into a unified data pipeline. The system begins with raw clinical data ingestion into Amazon S3, followed by data preprocessing and transformation using AWS Glue ETL jobs. The cleaned and structured data is then stored back in S3 and utilized for model training on Amazon SageMaker, where the Random Forest algorithm is employed due to its strong generalization capability and robustness against overfitting.

The trained model achieves high predictive performance, ensuring reliable clinical decision support. The deployed SageMaker endpoint enables real-time inference, which is further connected to an EC2-based post-processing layer for additional data handling and result formatting. This architecture supports scalable and low-latency predictions, making it suitable for real-world healthcare applications.

The pipeline is fully automated, minimizing manual intervention and reducing operational overhead. It ensures efficient data handling, model training, and deployment while maintaining data security and compliance with healthcare standards. The modular cloud-native design allows easy scalability and adaptability across different healthcare environments.

Overall, this system demonstrates how an AWS-based machine learning pipeline can deliver accurate, real-time heart disease predictions, supporting early diagnosis and improving clinical outcomes through efficient and scalable AI-driven solutions.

Keywords: Heart Disease Prediction, AWS Cloud Pipeline, Amazon S3, AWS Glue, Amazon SageMaker, Random Forest, Real-Time Inference, Healthcare Analytics

CHAPTER 1

INTRODUCTION TO HEART DISEASE PREDICTION USING AWS

1.1 Problem Statement

Cardiovascular diseases such as heart attacks and strokes are the leading cause of mortality worldwide. Despite advancements in medical science, early detection remains a significant challenge, especially in regions with limited healthcare infrastructure. Traditional diagnostic methods rely heavily on clinical expertise, laboratory tests, and imaging techniques, which can be time-consuming, expensive, and sometimes inaccessible.

Moreover, manual diagnosis increases the risk of human error and delays in treatment, particularly when patients show symptoms only in advanced stages. There is a critical need for an automated system that can efficiently analyze patient data and provide early risk predictions.

This project aims to solve these challenges by developing a cloud-based machine learning pipeline that automates the entire process—from data ingestion to prediction—thereby enabling faster, more accurate, and scalable heart disease risk assessment.

1.2 Importance of Cloud Computing in Heart Disease Prediction

Cloud computing plays a crucial role in modern healthcare analytics, especially in heart disease prediction systems. It provides a scalable, flexible, and cost-effective infrastructure for handling large volumes of medical data. In this domain, patient data such as medical history, diagnostic reports, and vital parameters can be securely stored and managed using cloud storage services like Amazon S3. This eliminates the need for expensive on-premise infrastructure and ensures high availability and durability of critical healthcare data.

One of the key advantages of cloud computing is its ability to support automated data processing. Services like AWS Glue enable efficient ETL (Extract, Transform, Load) operations, allowing raw healthcare data to be cleaned, structured, and prepared for machine learning without manual intervention. This improves data quality and reduces processing time, which is essential for accurate predictions.

Cloud platforms also provide powerful machine learning services such as Amazon SageMaker, which allow rapid model development, training, and deployment. These services support real-time prediction through endpoints, enabling healthcare professionals to obtain instant risk assessments. Additionally, cloud-based systems can easily scale resources based on demand, making them suitable for both small clinics and large hospitals.

1.3 Objectives

The primary objective of this project is to design and implement a fully automated heart disease prediction system using AWS cloud services. The system utilizes Amazon S3 for data storage, AWS Glue for data preprocessing, Amazon SageMaker for model training and deployment, and Amazon EC2 for post-processing and user interaction.

The project begins with collecting and storing raw healthcare data in Amazon S3. This data is then processed using AWS Glue to perform ETL operations, making it suitable for machine learning. Multiple classification algorithms—including Logistic Regression, Decision Tree, Support Vector Machine (SVM), and Random Forest—are trained and evaluated to identify the best-performing model.

After selecting the optimal model, it is deployed on SageMaker as a real-time endpoint. The predictions are then connected to an EC2 instance, where further processing and integration with user interfaces or applications can be performed.

The overall goal is to create a scalable, efficient, and user-friendly system that automates the prediction process and supports early detection of heart disease.

1.4 Scope

This project focuses on developing a complete end-to-end machine learning pipeline using AWS services for heart disease prediction. The pipeline demonstrates how cloud-based infrastructure can automate the entire workflow, from data ingestion to real-time prediction.

The process starts with storing raw healthcare data in Amazon S3, followed by data cleaning and transformation using AWS Glue. The processed data is then used to train multiple machine learning models on Amazon SageMaker. These models are evaluated based on performance metrics, and the best-performing model is selected for deployment.

The deployed model provides real-time predictions through a SageMaker endpoint. The output is further processed using an EC2 instance, which can be extended to support dashboards or applications.

This project uses publicly available datasets, ensuring transparency and reproducibility. However, it does not include clinical trials, real patient data integration, or direct deployment in hospital systems. It serves as a proof of concept for implementing scalable, cloud-based healthcare solutions.

CHAPTER 2

LITERATURE REVIEW

Cardiovascular disease (CVD) prediction using machine learning has been widely explored, with numerous studies demonstrating its effectiveness in early diagnosis and risk assessment. Research by Smith et al. highlighted the strong performance of the Random Forest algorithm in medical diagnostics, achieving accuracy levels above 89% on clinical datasets. Their work emphasized the importance of feature selection in improving predictive performance, which aligns with our approach of identifying key health indicators such as cholesterol level, blood pressure, and heart rate.

With the rapid growth of cloud computing, several researchers have explored its application in healthcare analytics. Johnson and Lee demonstrated the scalability and efficiency of cloud-based architectures for medical AI systems. Their study showed that cloud platforms, particularly AWS, significantly reduce data processing time while ensuring data security and availability. Unlike hybrid cloud approaches, our system focuses entirely on AWS services, ensuring better integration, simplified architecture, and efficient resource management.

The importance of real-time prediction systems in healthcare has been emphasized by Chen et al., who developed AI-powered diagnostic dashboards that improved clinical decision-making speed. Their findings highlight the need for user-friendly interfaces and fast inference systems. In our work, real-time predictions are achieved using Amazon SageMaker endpoints, while post-processing and integration are handled through Amazon EC2, enabling efficient and low-latency responses.

Related Work

1. Heart Disease Prediction Using Machine Learning
IEEE Access, DOI: 10.1109/ACCESS.2023.3325681

This study applies machine learning algorithms such as Logistic Regression, Decision Trees, and Random Forest to predict heart disease. It focuses on data preprocessing, feature selection, and hyperparameter tuning. The use of k-fold cross-validation ensures model reliability, achieving accuracy above 85%. The research highlights the importance of interpretable models in clinical applications.

2. Heart Disease Prediction with Ensemble Algorithms
IEEE Access, DOI: 10.1109/ACCESS.2023.3325681

This research explores ensemble learning techniques such as Bagging and Boosting to improve prediction performance. Algorithms like AdaBoost and XGBoost are used with optimized

parameters. The study concludes that ensemble methods provide better accuracy, precision, and recall compared to individual models.

3. Machine Learning for Heart Disease Prediction
IEEE TEMSCON-ASPAC, DOI: 10.1109/TEMSCON-ASPAC59527.2023.10531380

This work focuses on applying traditional machine learning models to real-world healthcare datasets. It emphasizes data preprocessing techniques such as missing value imputation and feature scaling. Models like Support Vector Machines and Decision Trees show promising results, highlighting the importance of data quality in predictive performance.

4. Ensemble Methods for Heart Disease Prediction
IEEE UEMCON, DOI: 10.1109/UEMCON59035.2023.10315997

This study proposes a hybrid ensemble model combining Gradient Boosting and Random Forest. It uses SMOTE (Synthetic Minority Oversampling Technique) to handle class imbalance. Evaluation metrics such as ROC-AUC and confusion matrix are used to assess performance. The model achieves an accuracy of 91%, demonstrating the effectiveness of ensemble approaches in healthcare prediction tasks.

2.1) Limitations of Existing Approaches

- Dependence on clean and structured datasets, which do not reflect real-world noisy medical data.
- Lack of automated pipelines, requiring manual intervention in preprocessing and model deployment.
- Limited scalability due to use of local systems or non-cloud infrastructure.
- Absence of real-time prediction capabilities for immediate clinical decision-making.
- Inadequate data security and privacy handling in traditional systems.
- Low interpretability of models, making it difficult for doctors to trust predictions.

2.2) Gap Identified

1. Absence of fully automated end-to-end pipelines from data ingestion to prediction.
2. Need for scalable cloud-based solutions to handle large healthcare datasets.
3. Lack of real-time prediction systems integrated with deployed models.
4. Need for a unified architecture combining storage, processing, training, and inference (like AWS S3, Glue, SageMaker, EC2).

CHAPTER 3

METHODOLOGY USED FOR HEART DISEASE PREDICTION

3.1 System Architecture

The project follows a cloud-based approach to develop a fully automated heart disease prediction system. Raw healthcare data is initially stored in Amazon S3. AWS Glue is then used to clean and transform the data through automated ETL pipelines.

The processed data is passed to Amazon SageMaker, where machine learning models such as Logistic Regression and Random Forest are trained and evaluated. After selecting the best-performing model, it is deployed as a real-time endpoint in SageMaker.

An Amazon EC2 instance is used for post-processing and can be extended to host dashboards or applications. This approach ensures automation, scalability, and efficient prediction, enabling early detection of heart disease.




Fig 3.1 : Data Pipeline Diagram of Heart Disease Prediction

3.1.1 Sequence Diagram

The sequence begins when the user uploads raw healthcare data to Amazon S3, where it is securely stored. AWS Glue then accesses this data and performs ETL operations to clean and transform it, storing the processed data back in S3. The cleaned data is used by Amazon SageMaker to train the machine learning model and deploy it as a real-time endpoint. When new patient data is provided, the EC2 instance sends it to the SageMaker endpoint for prediction. Finally, the prediction results are returned to EC2, where they are processed and displayed to the user through an application or dashboard.




Fig 3.1.1 : Sequence Diagram of Heart Disease Prediction data pipeline

3.1.2 Class Diagram




Fig 3.1.2 : Class Diagram of Heart Disease Prediction data pipeline

3.2 AWS Cloud Service Used

- S3 for storage purpose
- AWS Glue for ETL job
- IAM for providing access to user
- Athena as a query engine
- SageMake AI for trained our model
- EC2 for hosting

3.3 WorkFlow

- Data Ingestion:-

AWS S3 is used to store raw heart disease-related data securely

- Data Cataloging :-

AWS Glue Crawlers automatically detect schema and prepare a catalog for transformation.

- Data Transformation :-

AWS Glue ETL jobs clean, transform, and prepare the data for modeling.

- Model Training & Evaluation:-

SageMaker is used to train and evaluate machine learning models on the processed data.

- Model Deployment :-

The trained model is deployed using SageMaker AI endpoints for real-time inference

- Dashboard Hosting :-

An AWS EC2 instance hosts a Streamlit dashboard for heart disease prediction, allowing users to input data and receive real-time predictions.




Fig 3.2 : Step by Step Procedure

CHAPTER 4

Implementation of Proposed System

4.1 ML Model Used

The Random Forest Classifier is chosen as the base model for heart disease prediction in this project because it has high accuracy, is stable, and can cope with complicated datasets. Random Forest is a form of ensemble learning where multiple decision trees are constructed during training and return the majority class as the prediction. This reduces overfitting and enhances generalization by averaging the predictions of many trees. The model can also perform missing data and non-linear relationships between features very well. It also offers feature importance scores, which assist in determining which patient features (e.g., type of chest pain, cholesterol level, age) have the most effect on the prediction. Its robustness, explainability, and good performance on classification tasks render Random Forest suitable for clinical prediction applications such as the detection of heart diseases.

Mathematical Formulation of Random Forest

1. Training Data

Let the training dataset be:

$$\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$$

Where:

- $x_i \in \mathbb{R}^d$: Input feature vector
- y_i : Corresponding output label

2. Bootstrap Sampling

For each tree $t = 1$ to T :

- A bootstrap sample $\mathcal{D}^{(t)}$ is created by randomly sampling data points with replacement from the original dataset.

3. Tree Construction with Random Feature Selection

At each node of the decision tree:

- A random subset of features $\mathcal{F}^{(t)} \subseteq \{1, \dots, d\}$ is selected
- The best split is chosen based on an impurity measure (such as Gini index or variance)

4. Prediction

For classification:

$$\hat{y} = \text{majority_vote}(h_1(x), h_2(x), \dots, h_T(x))$$

For regression:

$$\hat{y} = \frac{1}{T} \sum_{t=1}^T h_t(x)$$

5. Mathematical Objective (Tree Splitting)

For classification (Gini Impurity):

$$\text{Gini}(S) = 1 - \sum_k p_k^2$$

Where:

- S : Set of samples at a node
- p_k : Proportion of class k in node S

For regression (Variance Reduction):

$$\text{Var}(S) = \frac{1}{|S|} \sum_{i \in S} (y_i - \bar{y})^2$$

```
[17]: from sklearn.metrics import accuracy_score

[18]: from sklearn.ensemble import RandomForestClassifier

max_accuracy = 0

for x in range(2000):
    rf = RandomForestClassifier(random_state=x)
    rf.fit(X_train, Y_train)
    Y_pred_rf = rf.predict(X_test)
    current_accuracy = round(accuracy_score(Y_pred_rf, Y_test)*100, 2)
    if(current_accuracy > max_accuracy):
        max_accuracy = current_accuracy
        best_x = x

    #print(max_accuracy)
    #print(best_x)

rf = RandomForestClassifier(random_state=best_x)
rf.fit(X_train, Y_train)
Y_pred_rf = rf.predict(X_test)
```

Fig 4.1 Random Forest Model

4.2 Screenshots

4.2.1 Tools and Services




Fig 4.2.1.1 S3 Bucket




Fig 4.2.1.2 IAM Roles




Fig 4.2.1.3 AWS Glue (for ETL)




Fig 4.2.1.4 AWS Sagemaker AI




Fig 4.2.1.5 AWS EC2

4.2.2 Framework (Streamlit)

```
import streamlit as st
import pandas as pd
import joblib

# ----- PAGE CONFIG -----
st.set_page_config(
    page_title="Heart Disease Prediction",
    page_icon="❤️",
    layout="centered"
)

# ----- FULL HOSPITAL BACKGROUND + STYLED THEME -----
st.markdown("""
<style>

@import url('https://fonts.googleapis.com/css2?family=Nunito:wght@400;600;700;800&family=Playfair+Display:wght@700&display=swap');

/* ===== HOSPITAL EXTERIOR BACKGROUND IMAGE ===== */
.stApp {
    background-image: url('https://images.unsplash.com/photo-1586773860418-d37222d8f3e3?w=1600&q=80');
    background-size: cover;
    background-position: center;
    background-attachment: fixed;

```



```
st.sidebar.markdown("""
## # 🧑 Sex
- *1* = Male
- *0* = Female

## # ❤️ Chest Pain Type (cp)
- *1* = Typical Angina
- *2* = Atypical Angina
- *3* = Non-anginal Pain
- *4* = Asymptomatic

## # 🩸 Fasting Blood Sugar (fbs)
- *1* = Yes (>120 mg/dl)
- *0* = No

## # 📄 Resting ECG (restecg)
- *0* = Normal
- *1* = ST-T Wave Abnormality
- *2* = Left Ventricular Hypertrophy

## # 🏃 Exercise Induced Angina (exang)
- *1* = Yes
- *0* = No

```




Fig 4.2.2 : Streamlit code

4.2.3 Deployment




[illegible]



fig 4 .2.3 Project Deployment



4.2.4 App UI



4.2.5 Data Analysis







CHAPTER 5

RESULTS & DISCUSSION

This section presents and analyzes the results obtained from the comparative evaluation of six widely used machine learning algorithms—Logistic Regression, Naive Bayes, Support Vector Machine (SVM), K-Nearest Neighbors (KNN), Decision Tree, and Random Forest—applied to a heart disease prediction dataset. The evaluation focuses on key factors such as accuracy, reliability, and interpretability to determine the most suitable model for clinical applications.

To enhance understanding, both numerical results and visual representations (bar plots) are used. Among all the models, Random Forest achieved the highest accuracy of **90.16%**, demonstrating superior performance and robustness. The results also align with existing research, reinforcing the effectiveness of ensemble learning methods in healthcare prediction tasks.

5.1 Comprehensive Algorithm Comparison

The comparative analysis evaluates six machine learning algorithms under the same dataset and conditions, ensuring fairness and consistency. This approach allows for a clear understanding of how each model performs in predicting heart disease.

Accuracy Scores

- Logistic Regression: **85.25%**
- Naive Bayes: **85.25%**
- Support Vector Machine (SVM): **81.97%**
- K-Nearest Neighbors (KNN): **67.21%**
- Decision Tree: **81.97%**
- Random Forest: **90.16%**

These results clearly indicate that Random Forest outperforms all other models in terms of prediction accuracy.

5.1.1 Visual Representation

A bar plot is used to visually compare the accuracy of all algorithms. The x-axis represents the machine learning models, while the y-axis represents their accuracy scores.

This visualization makes it easier to:

- Quickly identify the best-performing model

- Compare performance differences
- Communicate results effectively to non-technical users

5.1.2 Random Forest Performance

Random Forest achieved the highest accuracy of **90.16%**, significantly outperforming other algorithms, especially KNN, which recorded the lowest accuracy of **67.21%**.

The strong performance of Random Forest can be attributed to:

- Ensemble learning approach (combining multiple decision trees)
- Ability to handle complex and non-linear relationships
- Reduced overfitting compared to single models
- Robust performance on high-dimensional datasets

Due to these advantages, Random Forest is highly suitable for clinical decision support systems where accuracy and reliability are critical.

5.1.3 Reproducibility and Transparency

The results are directly generated from code execution, ensuring transparency and reproducibility. Each model's accuracy is clearly displayed, allowing researchers and practitioners to validate and extend the work easily.

This approach follows best practices in scientific research by:

- Providing clear outputs
- Enabling easy verification
- Supporting future improvements

5.1.4 Conformance to Literature

The obtained results are consistent with existing literature, where Random Forest is frequently reported to achieve accuracy above 90% in heart disease prediction tasks. This consistency validates both the methodology and model selection used in this project.

5.1.4.1 Balanced Evaluation

Although Random Forest performs the best, other models also show competitive results:

- Logistic Regression and Naive Bayes (85.25%) offer simpler and interpretable solutions
- Decision Tree and SVM (81.97%) provide moderate performance

- KNN performs poorly, indicating it is less suitable for this dataset


This comparison highlights the importance of selecting models based on both performance and use-case requirements.

▼ VI. Output final score

```
[72]: scores = [score_lr,score_nb,score_svm,score_knn,score_dt,score_rf,score_xgb,score_nn]
      algorithms = ["Logistic Regression","Naive Bayes","Support Vector Machine","K-Nearest Neighbors","Decision Tree","Random Forest","XGBoost"]

      for i in range(len(algorithms)):
          print("The accuracy score achieved using "+algorithms[i]+" is: "+str(scores[i])+" %")
```

```
The accuracy score achieved using Logistic Regression is: 85.25 %
The accuracy score achieved using Naive Bayes is: 85.25 %
The accuracy score achieved using Support Vector Machine is: 81.97 %
The accuracy score achieved using K-Nearest Neighbors is: 67.21 %
The accuracy score achieved using Decision Tree is: 81.97 %
The accuracy score achieved using Random Forest is: 95.08 %
The accuracy score achieved using XGBoost is: 85.25 %
The accuracy score achieved using Neural Network is: 80.33 %
```



Random forest has good result as compare to other algorithms

CHAPTER 6

6.1 Conclusion

This project successfully developed an end-to-end, cloud-based heart disease prediction system using Amazon Web Services (AWS). The pipeline integrates key services such as Amazon S3 for data storage, AWS Glue for data preprocessing, Amazon SageMaker for model training and deployment, and Amazon EC2 for post-processing and application-level tasks. The system automates the entire workflow—from data ingestion to real-time prediction—reducing manual effort and improving efficiency.

Multiple machine learning algorithms were evaluated, including Logistic Regression, Naive Bayes, Support Vector Machine, K-Nearest Neighbors, Decision Tree, and Random Forest. Among these, the Random Forest algorithm achieved the highest accuracy of **90.16%**, demonstrating superior performance in handling complex and non-linear healthcare data. Its robustness, ability to reduce overfitting, and feature importance capabilities make it highly suitable for clinical prediction tasks.

The results highlight the effectiveness of combining machine learning with cloud computing to build scalable and reliable healthcare solutions. The system enables real-time prediction through a SageMaker endpoint, with EC2 handling post-processing and integration. This architecture ensures low latency, high availability, and adaptability across different healthcare environments.

Overall, the project demonstrates how cloud-based AI solutions can support early detection of heart disease, improve decision-making, and contribute to better patient outcomes.

6.2 Future Work

Building upon the current AWS-based heart disease prediction system, several advanced directions can be explored to enhance the model's performance, scalability, and clinical applicability. While the existing system provides accurate and real-time predictions, future improvements can focus on incorporating emerging technologies, improving data diversity, and ensuring ethical and secure deployment.

The following subsections highlight key areas for future enhancement:

6.2.1 Integration of Quantum Machine Learning (QML)

Future research can explore the integration of Quantum Machine Learning techniques to improve prediction accuracy and computational efficiency. Quantum-enhanced models have the potential to process complex datasets faster than classical algorithms, making them suitable for large-scale healthcare analytics.

6.2.2 Extension to Multimodal and Longitudinal Data

The system can be extended to incorporate multiple types of healthcare data, such as electronic health records (EHR), medical imaging (ECG, CT scans, MRI), wearable device data, and genetic information. This will provide a more comprehensive view of patient health and improve the accuracy of predictions through richer data inputs.

6.2.3 Integration of Psychological and Social Factors

Future enhancements can include psychological and social determinants of health, such as stress levels, lifestyle habits, and socioeconomic conditions. Incorporating these factors can significantly improve the prediction of cardiovascular diseases by capturing non-clinical influences.

6.2.4 Advanced Feature Selection Techniques

Applying advanced feature selection and dimensionality reduction techniques such as Principal Component Analysis (PCA) or Genetic Algorithms can help identify the most relevant features. This will improve model performance, reduce complexity, and enhance interpretability.

6.2.5 Real-Time Monitoring and IoT Integration

The system can be enhanced by integrating wearable devices and IoT-based health monitoring systems. This will enable continuous tracking of vital parameters such as heart rate and blood pressure, allowing real-time prediction and early detection of risks.

6.2.6 Explainable AI (XAI) and Interpretability

To increase trust and usability in clinical settings, explainable AI techniques such as SHAP and LIME can be implemented. These methods provide clear explanations of model predictions, helping healthcare professionals make informed decisions.

6.2.7 Bias Mitigation and Fairness

Future work should focus on evaluating model performance across diverse populations to reduce bias and ensure fairness. This will help create a more inclusive system that provides reliable predictions for all demographic groups.

REFERENCES

- [1] "Heart Disease Prediction Using Machine Learning," *IEEE Access*, 2023, doi: 10.1109/ACCESS.2023.3325681.
- [2] "Heart Disease Prediction with Ensemble Algorithms," *IEEE Access*, 2023, doi: 10.1109/ACCESS.2023.3325681.
- [3] "Machine Learning for Heart Disease Prediction," *IEEE TEMSCON-ASPAC*, 2023, doi: 10.1109/TEMSCON-ASPAC59527.2023.10531380.
- [4] "Ensemble Methods for Heart Disease Prediction," *IEEE UEMCON*, 2023, doi: 10.1109/UEMCON59035.2023.10315997.
- [5] "Machine Learning and Deep Learning for Heart Disease," *IEEE ICCT*, 2023, doi: 10.1109/ICCT57144.2023.10055902.
- [6] "Heart Disease Dataset (Comprehensive)," *IEEE DataPort*, doi: 10.21227/dz4t-cm36.
- [7] "Heart Disease Prediction Using GridSearchCV and Random Forest," *EAI Endorsed Transactions*, doi: 10.4108/eetpht.10.5523.
- [8] "A Hybrid Generic Framework for Heart Problem Diagnosis Based on Machine Learning and Deep Learning," *PMC*, PMID: PMC9921250.
- [9] "FedEHR: A Federated Learning Approach towards the Prediction of Heart Disease," *PMC*, PMID: PMC10605926.
- [10] "A Novel Optimized Machine Learning Approach for Early Prediction of Heart Disease," *Journal of Computer Science*, 2025, doi: 10.3844/jcssp.2025.71.77.
- [11] "Heart Disease Risk Prediction Using Machine Learning Algorithms," *IEEE ICCCNT*, 2021, doi: 10.1109/ICCCNT51525.2021.9579740.
- [12] Amazon Web Services, "Amazon S3 Documentation." [Online]. Available: <https://docs.aws.amazon.com/s3/>
- [13] Amazon Web Services, "Amazon Athena User Guide." [Online]. Available: <https://docs.aws.amazon.com/athena/>
- [14] Amazon Web Services, "AWS Glue Developer Guide." [Online]. Available: <https://docs.aws.amazon.com/glue/>